

SDA Vision / Research report

infographic.pdf

Global Intelligence Summary

Information Epidemiology

Tracking the Digital Contagion of
Maternal Vaccine Misinformation

Synthesised from a 6-month, 15-country social media surveillance operation.



Assessment summary

Field	Recorded detail
Finding	Inconclusive
Recorded verdict	inconclusive
Confidence	low
Indicative rating	26/100
Rating basis	rounded mean of available frame scores; not a calibrated probability
Explanation	The sampled frames leave the assessment inconclusive. Review the frame results and source context.
Coverage	pdf - 4 frame(s) analysed of 15 found
Review context	Research assessment. Review alongside source information and context. When credentials are missing, origin remains unresolved by this check. Use these findings to inform a documented human review.

Provider overview

Provider	Type / status	Assessment	Rating
Claude Vision	vision / ok	inconclusive	55/100
OpenAI Vision	vision / ok	inconclusive	15/100
Gemini Vision	vision / ok	authentic likely	28/100
C2PA Content Credentials	provenance / ok	inconclusive	n/a
Local forensic cues	forensic / ok	partially synthetic	41/100

Provider evidence / 1

Field	Recorded detail
Provider	Claude Vision
Model	claude-opus-5-5
Status	ok
Assessment	inconclusive
Confidence	low
Role in assessment	Additional finding
Provider summary	Worst of 4 frames (figure p15): synthetic likelihood 55/100; tells: no camera-photo content; digitally rendered slide graphic, so photographic forensics (grain, lens blur, lighting) do not apply, irregular, non-geometric low-poly line networks with inconsistent node joins and fading, typical of AI image-generator decorative motifs, red-to-green left/right mirrored decoration with loosely matched structure, heavy bold text with slightly soft drop-shadow/halo edges on headline
Evidence 1	Frame figure p1: 45/100
Evidence 2	Frame figure p6: 40/100
Evidence 3	Frame figure p10: 35/100
Evidence 4	Frame figure p15: 55/100
Evidence 5	No camera-photo content; digitally rendered slide graphic, so photographic forensics (grain, lens blur, lighting) do not apply
Evidence 6	Irregular, non-geometric low-poly line networks with inconsistent node joins and fading, typical of AI image-generator decorative motifs
Evidence 7	Red-to-green left/right mirrored decoration with loosely matched structure
Evidence 8	Heavy bold text with slightly soft drop-shadow/halo edges on headline
Evidence 9	Slightly uneven centred line breaks in body text
Evidence 10	All text is correctly spelled and coherent, which is consistent with either modern text-capable generators or conventional slide software
Evidence 11	Faint mottled background texture rather than flat vector fill
Evidence 12	Caveat: This is a text-and-graphic slide, likely extracted from a PDF, not a photograph. Slides made in presentation software and slides from text-capable AI image generators can look alike. The irregular network motifs and soft rendering suggest possible AI generation, but they are not decisive. A human should check the source PDF's metadata, fonts and vector layers, and the authoring tool.

Provider evidence / 2

Field	Recorded detail
Provider	OpenAI Vision
Model	gpt-6-astra
Status	ok
Assessment	inconclusive
Confidence	low
Role in assessment	Additional finding
Provider summary	Worst of 4 frames (figure p6): synthetic likelihood 15/100; tells: Consistent, legible typography throughout both columns, Regular rounded panels and aligned borders characteristic of conventional slide design, No obvious malformed lettering or incoherent graphic elements, Entire image is a text infographic, with no photographic detail to assess
Evidence 1	Frame figure p1: 8/100
Evidence 2	Frame figure p6: 15/100
Evidence 3	Frame figure p10: 5/100
Evidence 4	Frame figure p15: 15/100
Evidence 5	Consistent, legible typography throughout both columns
Evidence 6	Regular rounded panels and aligned borders characteristic of conventional slide design
Evidence 7	No obvious malformed lettering or incoherent graphic elements
Evidence 8	Entire image is a text infographic, with no photographic detail to assess
Evidence 9	Caveat: A conventional infographic and an AI-assisted graphic can look identical; source files or creation history are needed to establish provenance.

Provider evidence / 3

Field	Recorded detail
Provider	Gemini Vision
Model	gemini-3.8-flash
Status	ok
Assessment	authentic likely
Confidence	medium
Role in assessment	Additional finding
Provider summary	Worst of 4 frames (figure p15): synthetic likelihood 28/100; tells: clean vector typography, standard geometric graphic design elements, absence of diffusion generation artifacts
Evidence 1	Frame figure p1: 25/100
Evidence 2	Frame figure p6: 25/100
Evidence 3	Frame figure p10: 25/100
Evidence 4	Frame figure p15: 28/100
Evidence 5	Clean vector typography
Evidence 6	Standard geometric graphic design elements
Evidence 7	Absence of diffusion generation artifacts
Evidence 8	Caveat: Standard presentation slide with digital vector art; non-photographic graphics make AI detection inconclusive.

Content Credentials / provenance

Field	Recorded detail
Provider	C2PA Content Credentials
Model	Not recorded
Status	ok
Assessment	inconclusive
Confidence	Not recorded
Role in assessment	Additional finding
Provider summary	Embedded credentials: none found. Origin remains unresolved by this check.
Evidence 1	Coverage: credentials embedded in the submitted file. External credentials were not retrieved.
Evidence 2	Verifier: c2pa-python 0.37.12; native SDK 0.91.0; specification target 2.4; policy embedded-offline-reviewed-trust-v3.
Evidence 3	Offline embedded-manifest verification; no network trust-list, manifest or revocation refresh.
Evidence 4	No reviewed local C2PA trust snapshot installed. Provenance overrides are disabled.

Provider evidence / 5

Field	Recorded detail
Provider	Local forensic cues
Model	Not recorded
Status	ok
Assessment	partially synthetic
Confidence	low
Role in assessment	Additional finding
Provider summary	Worst of 4 frames (figure p1): Heuristic synthetic likelihood 41/100 from metadata + pixel statistics.
Evidence 1	Frame figure p1: 41/100
Evidence 2	Frame figure p6: 41/100
Evidence 3	Frame figure p10: 41/100
Evidence 4	Frame figure p15: 41/100
Evidence 5	No camera EXIF (common in AI exports, screenshots, and re-saves).
Evidence 6	ELA mean 0.4, peak 23 (uniform high ELA can indicate generation/heavy edit).
Evidence 7	High-frequency noise std 14.12 (very low/very uniform noise is a synthesis cue).
Evidence 8	Spectral peakiness 19.47 (periodic frequency peaks can mark generated textures).

Per-frame results

Frame	Assessment	Rating / 100
figure p1	authentic likely	25
figure p6	inconclusive	25
figure p10	authentic likely	25
figure p15	inconclusive	28

Evidence supporting this verdict

Field	Recorded detail
Record	No entries recorded.

Additional findings

Field	Recorded detail
Claude Vision	Worst of 4 frames (figure p15): synthetic likelihood 55/100; tells: no camera-photo content; digitally rendered slide graphic, so photographic forensics (grain, lens blur, lighting) do not apply, irregular, non-geometric low-poly line networks with inconsistent node joins and fading, typical of AI image-generator decorative motifs, red-to-green left/right mirrored decoration with loosely matched structure, heavy bold text with slightly soft drop-shadow/halo edges on headline
OpenAI Vision	Worst of 4 frames (figure p6): synthetic likelihood 15/100; tells: Consistent, legible typography throughout both columns, Regular rounded panels and aligned borders characteristic of conventional slide design, No obvious malformed lettering or incoherent graphic elements, Entire image is a text infographic, with no photographic detail to assess
Gemini Vision	Worst of 4 frames (figure p15): synthetic likelihood 28/100; tells: clean vector typography, standard geometric graphic design elements, absence of diffusion generation artifacts
Local forensic cues	Supporting observation; considered alongside the model assessments.
C2PA Content Credentials	Embedded credentials: none found. Origin remains unresolved by this check.

Extracted text

Field	Recorded detail
Recorded text	Geographic Discourse Extremes The Sceptics (USA & Italy) • Primary Driver: Mistrust in institutions (CDC, FDA, Italy's ASL). • Core Theme: Extrapolating harm from contraindicated vaccines to recommended ones. • Vector Activity: High presence of bot-associated hashtags. The Advocates (India & UK) • Primary Driver: Proactive, visible public health campaigns. • Core Theme: Child protection and successful programme launches (e.g., India's Mission Indradhanush). • Vector Activity: Organic tagging of trusted bodies (Public Health England).

How this assessment was reached

Field	Recorded detail
1. frame_agreement	frames=4; synthetic_frames=0; authentic_frames=2; rating_spread=3
2. verdict	value=inconclusive; confidence=low; reason=The sampled frames leave the assessment inconclusive. Review the frame results and source context.

File and analysis record

Field	Recorded detail
File	infographic.pdf
Analysis ID	279c5518-74f1-42db-a265-19a0d0fcb0bc
Generated	2026-09-26T17:54:48+00:00
Tool	SDA Vision 0.1.0
Models	claude-opus-5-5, gemini-3.8-flash, gpt-6-astra
Method version	verdict-consistency-2026-09-26.1
Analysis lens	general
Original SHA-256	02915f051b6eb2c9ef0c21f5bb20c4229ca21324bf20172f3867c32c299a5924
Ingestion note 1	Extracted 4 embedded figure(s) of 15 found (figure mode).
Ingestion note 2	No selectable text layer - vision models read the page/figure images.
Ingestion note 3	Recomputed locally from the same saved provider replies acquired under run five-examples-20260926-r1. No new provider requests. Original report 67783866-5136-47c8-95ce-ff7e7f37c6d3 was generated at 2026-09-26T17:38:42+00:00 under method publication-repair-2026-09-24.3; that report is preserved.
Ingestion note 4	Method verdict-consistency-2026-09-26.1: directional model votes require matching explicit verdicts and rating thresholds. Inconclusive replies retain their ratings and remain undecided observations.
Ingestion note 5	Presentation copy of saved report aac43583-5a80-4042-988f-028d0b7bfae8. The descriptive filename replaces the earlier example ID for navigation. Original analysis dates, provider replies, scores and acquisition context are retained; no new model analysis.