

SDA Vision / Research report

podcast.m4a



Assessment summary

Field	Recorded detail
Finding	Leaning synthetic
Recorded verdict	synthetic likely
Confidence	high
Indicative rating	90/100
Rating basis	median of available model ratings; not a calibrated probability
Explanation	3 models agree this is likely synthetic.
Coverage	audio - 0 frame(s) analysed
Review context	Research assessment. Review alongside source information and context. When credentials are missing, origin remains unresolved by this check. Use these findings to inform a documented human review.

Provider overview

Provider	Type / status	Assessment	Rating
Claude Analysis	analysis / ok	synthetic likely	78/100
OpenAI Analysis	analysis / ok	synthetic likely	90/100
Gemini Analysis	analysis / ok	synthetic likely	95/100
Local forensic cues	forensic / ok	not applicable	n/a
C2PA Content Credentials	provenance / ok	inconclusive	n/a

Provider evidence / 1

Field	Recorded detail
Provider	Claude Analysis
Model	claude-opus-5-5
Status	ok
Assessment	synthetic likely
Confidence	medium
Role in assessment	Supporting
Provider summary	Style closely matches AI-generated two-host 'deep dive' audio overviews (e.g. NotebookLM): stock openers ('Imagine holding...', 'Okay, let's unpack this', 'If we connect this to the bigger picture'), constant affirmations ('Exactly', 'Perfectly said'), scripted neutrality disclaimer and a staged host challenge, suggesting likely synthetic origin.
Evidence 1	Synthetic origin framing: Two-host explainer podcast summarising a peer-reviewed study, with historical context, a stated neutrality disclaimer and interpretive commentary on polarisation
Evidence 2	Summary (content, for the researcher's coding):
Evidence 3	Study by Sam Martin and Samantha Vanderslott (Oxford Vaccine Group) published in journal Vaccine
Evidence 4	Analysed 7.89 million English-language US/UK tweets, June 2020 to June 2021
Evidence 5	Used Meltwater for collection and InfraNodus for text network analysis
Evidence 6	5.91 million tweets on masks, 1.98 million on vaccines
Evidence 7	Caveat: Human podcasters can use similar conventions, and transcription artefacts may mask or mimic synthetic patterns; source audio and provenance should be checked.

Provider evidence / 2

Field	Recorded detail
Provider	OpenAI Analysis
Model	gpt-6-astra
Status	ok
Assessment	synthetic likely
Confidence	medium
Role in assessment	Supporting
Provider summary	The highly regular two-host exchanges, repetitive affirmations, scripted explanatory prompts, and formulaic transitions strongly suggest synthetic generation or editing.
Evidence 1	Synthetic origin framing: Conversational research explainer using staged questions, emphatic agreement, historical parallels, and sociological interpretations.
Evidence 2	Summary (content, for the researcher's coding):
Evidence 3	The cited study reportedly analyzed 7.89 million English-language tweets from June 2020 to June 2021.
Evidence 4	Researchers reportedly used Meltwater and InfraNodus to examine mask and vaccine discourse in the US and UK.
Evidence 5	The text describes convergence between anti-vaccine and anti-mask groups.
Evidence 6	It reports greater negative mask sentiment in the UK than the US.
Evidence 7	Caveat: A human-scripted podcast could share these features; transcript artifacts and absent provenance limit attribution.

Provider evidence / 3

Field	Recorded detail
Provider	Gemini Analysis
Model	gemini-3.8-flash
Status	ok
Assessment	synthetic likely
Confidence	high
Role in assessment	Supporting
Provider summary	The text displays the classic conversational tropes, rapid turn-taking, and synthetic cadence characteristic of an automated AI podcast generator like NotebookLM.
Evidence 1	Synthetic origin framing: An informal two-host podcast script discussing academic research on social media discourse and health mandates.
Evidence 2	Summary (content, for the researcher's coding):
Evidence 3	Oxford Vaccine Group researchers analyzed 7.89 million English-language tweets from the US and UK between June 2020 and June 2021
Evidence 4	Study used Meltwater and InfraNodus network analysis software to examine sentiment around masks and vaccines
Evidence 5	5.91 million analyzed tweets focused on masks, while 1.98 million focused on vaccines
Evidence 6	Negative sentiment toward masks was 28% in the UK compared to 22% in the US
Evidence 7	Caveat: Human podcast hosts discussing academic papers can occasionally use similar rehearsed, highly structured banter.

Provider evidence / 4

Field	Recorded detail
Provider	Local forensic cues
Model	Not recorded
Status	ok
Assessment	not applicable
Confidence	Not recorded
Role in assessment	Not listed
Provider summary	Pixel forensics: outside text analysis.

Content Credentials / provenance

Field	Recorded detail
Provider	C2PA Content Credentials
Model	Not recorded
Status	ok
Assessment	inconclusive
Confidence	Not recorded
Role in assessment	Additional finding
Provider summary	Embedded credentials: none found. Origin remains unresolved by this check.
Evidence 1	Coverage: credentials embedded in the submitted file. External credentials were not retrieved.
Evidence 2	Verifier: c2pa-python 0.37.12; native SDK 0.91.0; specification target 2.4; policy embedded-offline-reviewed-trust-v3.
Evidence 3	Offline embedded-manifest verification; no network trust-list, manifest or revocation refresh.
Evidence 4	No reviewed local C2PA trust snapshot installed. Provenance overrides are disabled.

Local soundtrack inspection

Field	Recorded detail
Record 1	Audio inspection: present; tracks: 1.
Record 2	Track 0: aac, 2 channel(s), 44100 Hz; duration 1468.569252s; decode decoded; 3/3 sampled windows above -60 dBFS after mono downmix.
Record 3	Sample windows: 0.0s + 3.0s; 732.785s + 3.0s; 1465.569s + 3.0s
Record 4	Original-file credentials are recorded separately. Credential coverage of this audio track: unresolved.
Record 5	Audio-level measurements cover the sampled windows. The remaining audio has not been inspected by this check.

Optional audio excerpt description

Field	Recorded detail
Record 1	Google Gemini / gemini-3.8-flash; status: ok; content: speech.
Record 2	Checked: 2026-09-27T19:20:45.393134+00:00; consent to Google recorded: True.
Record 3	Track 0; excerpt 0s + 20.0s; prompt: soundtrack-content-v1.
Record 4	Original SHA-256: 58663a6a9e4b634bdb703ead0aaa7cd6dce7fbf97e70c3c1cc61b3311b94c0b4; excerpt SHA-256: f9713796490ddc9624cfc92412738eaca7192d9cbde08dd259237e29e2aa2128.
Record 5	A female voice speaks in an interview or conversational tone.
Record 6	A male voice responds briefly in agreement.
Record 7	The conversation sounds like a podcast or radio broadcast discussion.
Record 8	Audio is clean with minimal to no background noise.
Record 9	Transcript excerpt: Imagine holding, uh, like a simple piece of cloth in your hand. Let's say it's 2019, and it's just a standard surgical mask. Right, just a boring everyday medical supply, something you'd really only see at the dentist's office. Exactly.
Record 10	This assessment describes the selected audio excerpt. Origin and credentials are considered separately.
Record 11	Analysed excerpt: mono, 16 kHz PCM. Credential checks use the original file.
Record 12	Audio findings are shown alongside the visual assessment and remain separate from its score.

Assessment coverage

Field	Recorded detail
Record 1	Visual verdicts for videos cover the sampled still frames. Standalone-audio content assessments cover the transcript. Soundtrack-origin detection and continuous-motion analysis are outside these assessments.
Record 2	Original-file credentials, excerpt descriptions and visual interpretation are recorded separately. Credential coverage of an individual audio track remains unresolved by this workflow.

Evidence supporting this verdict

Field	Recorded detail
Claude Analysis	Style closely matches AI-generated two-host 'deep dive' audio overviews (e.g. NotebookLM): stock openers ('Imagine holding...', 'Okay, let's unpack this', 'If we connect this to the bigger picture'), constant affirmations ('Exactly', 'Perfectly said'), scripted neutrality disclaimer and a staged host challenge, suggesting likely synthetic origin.
OpenAI Analysis	The highly regular two-host exchanges, repetitive affirmations, scripted explanatory prompts, and formulaic transitions strongly suggest synthetic generation or editing.
Gemini Analysis	The text displays the classic conversational tropes, rapid turn-taking, and synthetic cadence characteristic of an automated AI podcast generator like NotebookLM.

Additional findings

Field	Recorded detail
C2PA Content Credentials	Not checked in this text-analysis pass; no original-media provenance conclusion.

Transcribed text

Field	Recorded detail
Recorded text	Imagine holding a like a simple piece of cloth in your hand. Let's say it's 2019, and it's just a standard surgical mask. Right, just a boring everyday medical supply, something you'd really only see at the dentist office. Exactly. But then fast-forward to the end of 2020, and that exact same piece of cloth had morphed into like a literal measure of your morality. Oh, absolutely. It became a symbol of your political affiliation, your patriotism, even. It's just wild to think about. I mean, today we are looking at the exact digital moment a basic health mandate turned into a full-blown ideological war. It really is one of the most compressed, intense sociological shifts we've ever witnessed. Yeah. We went from a public that largely didn't think about epidemiology at all to a society where every single person was forced to take a public, highly visible stance on a global health crisis. And usually, you know, when we talk about a massive global event like that, we have to rely on historians decades later to piece together the public mood. Right, digging through old newspaper clippings or a diary entries. Exactly. But for the COVID-19 pandemic, we have this complete, real-time behavioral laboratory. So, we are grounding our conversation today in this fascinating peer-reviewed study from the journal Vaccine. Yes, authored by Sam Martin and Samantha Vanderslott from the Oxford Vaccine Group. And the scale of their research here is what makes it so robust, right? Oh, totally. They analyzed this massive data set of 7.89 million English-language tweets. Wow. Yeah, from the US and the UK, spanning a very specific window, basically June 1st, 2020 to June 1st, 2021. And they didn't just like count how many times the word mask or vaccine was used, which is what I would have assumed. No, they went way deeper. They utilized media monitoring software called Meltwater and text network analysis software called InfraNodus. Which uh we should probably explain how those tools work because it's vital for understanding the whole study. Please do. So Meltwater acts like this giant net scraping the internet and pulling in every public conversation based on specific keywords. Right. But InfraNodus acts as the brain. It actually maps relationships between words. So if the word mask is constantly sitting right next to the word tyranny or sheep in millions of tweets, The software draws a thick digital line between them. Exactly. It lets researchers see the exact architecture of public sentiment. It's almost like looking at a real-time MRI of society's collective anxiety. That's a really great way to visualize it. It moves us past just individual anecdotes, Yeah. Yeah. and it shows us the actual structural flow of ideas. Yeah. So, the mission for our deep dive today is to explore a very specific phenomenon uncovered in this flow. We want to see how pre-existing anti-vaccine groups and newly formed anti-mask groups actually merged into this massive digital powerhouse. And to understand the psychological, social, and political drivers behind how the public reacted to these mandates. Exactly. Now uh a very quick heads up to you the listener before we dive into the timeline, the data we are analyzing today is heavily political. Very heavily. Yeah, it involves major figures from the left and the right. We're talking President Joe Biden, former President Donald Trump, Vice President Kamala Harris along with a bunch of conservative and liberal viewpoints. Right. So, I have to be clear here. We aren't taking sides today. We aren't endorsing any of the politics discussed. We're just impartially reading the digital footprint left behind to help you understand the mechanics of this social media evolution. Our focus is strictly on the mechanics of the discourse. You know, what the data reflects, what was said, and how it spread. Okay, let's unpack this. To really understand the explosion of outrage and compliance in 2020, we have to recognize that human pushback again

How this assessment was reached

Field	Recorded detail
1. override	outcome=none; detail=No confirmed C2PA AI assertion or SynthID watermark. Origin remains unresolved by these checks.
2. model_agreement	synthetic=['Claude Analysis', 'OpenAI Analysis', 'Gemini Analysis']; authentic=[]; uncertain=[]; forensic_supports_synthetic=False
3. verdict	value=synthetic_likely; confidence=high; reason=3 models agree this is likely synthetic.

Cite this analysis

Field	Recorded detail
Software (Harvard)	Martin, S. (2026) SDA Vision: Synthetic-media Discourse Analysis (version 0.1.0) [Computer software]. Manchester Metropolitan University. Available at: https://sdavision.io/
Software creator	Dr Sam Martin, Smart Data Research UK (UKRI) Fellow, Manchester Metropolitan University (MMU). ORCID: https://orcid.org/0000-0002-4466-8374 .
Funding acknowledgement	This work was supported by Smart Data Research UK, a UKRI investment; Grant number UKRI4010.
Analysis record	File: podcast.m4a; media type: audio; recorded at: 2026-09-26T22:31:07+00:00; models: claude-opus-5-5, gemini-3.8-flash, gpt-6-astra.

File and analysis record

Field	Recorded detail
File	podcast.m4a
Analysis ID	7e1a6917-52a1-453d-8dda-9c4842232f31
Generated	2026-09-26T22:31:07+00:00
Tool	SDA Vision 0.1.0
Models	claude-opus-5-5, gemini-3.8-flash, gpt-6-astra
Method version	pdf-preparation-2026-09-26.1
Analysis lens	general
Original SHA-256	58663a6a9e4b634bdb703ead0aaa7cd6dce7fbf97e70c3c1cc61b3311b94c0b4
Ingestion note 1	Transcoded to 5737 KB mono 16 kHz (1800s).
Ingestion note 2	Transcribed with Gemini.
Ingestion note 3	Standalone-audio content assessment covers the transcript. AI-origin detection from sound is outside this assessment. Original-file C2PA was checked before transcription and is reported separately.
Ingestion note 4	Audio tracks inspected locally. Soundtrack uploads require a separate optional check; AI-origin detection from sound is outside this workflow.
Ingestion note 5	Fresh provider replies acquired for eleven-examples-refresh-20260926-r1; report assembled locally from the saved responses without further API calls.
Ingestion note 6	Transcript analysis used the first 12000 characters of a fresh 25985-character transcription; this report retains its first 4000 characters. Sound-origin detection is outside this check.
Ingestion note 7	Provider latency values are restored from the acquisition receipts (maximum sampled-frame latency per provider); raw replies and assessment results are unchanged by this bookkeeping step.